

Task 3

Task Description:

Fine-Tune a Language Model for Domain-Specific Customer Support

Project Title

Fine-Tuning an Open-Source LLM for Medical Conversations

Objectives

1. You will fine-tune a pre-trained open-source language model (such as a small LLaMA-based model from Meta ecosystem or a Hugging Face model from Hugging Face) to specialize in answering technical support questions.
2. The goal is to transform a general-purpose model into a domain-specific assistant.

Requirements

1. Load a small pre-trained model
2. Prepare dataset in instruction format
3. Apply a fine-tuning method: Preferably LoRA (Parameter-Efficient Fine-Tuning)
4. Train for a few epochs
5. Save the fine-tuned adapter
6. Compare responses: before fine-tuning and after fine-tuning